

Naïve Bayes Classifier

Eamonn Keogh
UCR



Thomas Bayes




1702 - 1761

This is a high level overview only. For details, see:
Pattern Recognition and Machine Learning, Christopher Bishop, Springer-Verlag, 2006.
Or
Pattern Classification by R. O. Duda, P. E. Hart, D. Stork, Wiley and Sons.

We will start off with a visual intuition, before looking at the math...


Grasshoppers

Katydids




Remember this example?
Let's get lots more data...


With a lot of data, we can build a histogram. Let us just build one for “Antenna Length” for now...



We can leave the histograms as they are, or we can summarize them with two normal distributions.




Let us use two normal distributions for ease of visualization in the following slides...



- We want to classify an insect we have found. Its antennae are 3 units long. How can we classify it?
- We can just ask ourselves, give the distributions of antennae lengths we have seen, is it more *probable* that our insect is a **Grasshopper** or a **Katydid**.
- There is a formal way to discuss the most *probable* classification...



$p(c_j | d)$ = probability of class c_j , given that we have observed d



$p(c_j | d)$ = probability of class c_j , given that we have observed d

$$P(\text{Grasshopper} | 3) = 10 / (10 + 2) = 0.833$$

$$P(\text{Katydid} | 3) = 2 / (10 + 2) = 0.166$$




Antennae length is 3


$p(c_j | d)$ = probability of class c_j , given that we have observed d

$$P(\text{Grasshopper} | 7) = 3 / (3 + 9) = 0.250$$

$$P(\text{Katydid} | 7) = 9 / (3 + 9) = 0.750$$




Antennae length is 7



$p(c_j | d)$ = probability of class c_j , given that we have observed d

$$P(\text{Grasshopper} | 5) = 6 / (6 + 6) = 0.500$$

$$P(\text{Katydid} | 5) = 6 / (6 + 6) = 0.500$$



Bayes Classifiers

That was a visual intuition for a simple case of the Bayes classifier, also called:

- Idiot Bayes
- Naïve Bayes
- Simple Bayes

We are about to see some of the mathematical formalisms, and more examples, but keep in mind the basic idea.

*Find out the probability of the **previously unseen instance** belonging to each class, then simply pick the most probable class.*

Bayes Classifiers

- Bayesian classifiers use **Bayes theorem**, which says

$$p(c_j | d) = \frac{p(d | c_j) p(c_j)}{p(d)}$$

- $p(c_j | d)$ = probability of instance d being in class c_j ,
This is what we are trying to compute
- $p(d | c_j)$ = probability of generating instance d given class c_j ,
We can imagine that being in class c_j , causes you to have feature d with some probability
- $p(c_j)$ = probability of occurrence of class c_j ,
This is just how frequent the class c_j , is in our database
- $p(d)$ = probability of instance d occurring

This can actually be ignored, since it is the same for all classes

Assume that we have two classes

$c_1 = \text{male}$, and $c_2 = \text{female}$.

We have a person whose sex we do not know, say “*drew*” or *d*.

Classifying *drew* as male or female is equivalent to asking is it more probable that *drew* is **male** or **female**, I.e which is greater $p(\text{male} \mid \text{drew})$ or $p(\text{female} \mid \text{drew})$

(Note: “Drew can be a male or female name”)



Drew Barrymore



Drew Carey

What is the probability of being called “*drew*” given that you are a **male**?

What is the probability of being a **male**?

What is the probability of being named “*drew*”?

(actually irrelevant, since it is that same for all classes)

$$p(\text{male} \mid \text{drew}) = \frac{p(\text{drew} \mid \text{male}) p(\text{male})}{p(\text{drew})}$$



Officer Drew

This is Officer Drew (who arrested me in 1997). Is Officer Drew a **Male** or **Female**?

Luckily, we have a small database with names and sex.

We can use it to apply Bayes rule...

Name	Sex
Drew	Male
Claudia	Female
Drew	Female
Drew	Female
Alberto	Male
Karin	Female
Nina	Female
Sergio	Male

$$p(c_j | d) = \frac{p(d | c_j) p(c_j)}{p(d)}$$



Officer Drew

$$p(c_j | d) = \frac{p(d | c_j) p(c_j)}{p(d)}$$

Name	Sex
Drew	Male
Claudia	Female
Drew	Female
Drew	Female
Alberto	Male
Karin	Female
Nina	Female
Sergio	Male

$$p(\text{male} | \text{drew}) = \frac{1/3 * 3/8}{3/8} = \frac{0.125}{3/8}$$

$$p(\text{female} | \text{drew}) = \frac{2/5 * 5/8}{3/8} = \frac{0.250}{3/8}$$

Officer Drew is more likely to be a **Female**.



Officer Drew IS a female!

Officer Drew

$$p(\text{male} \mid \text{drew}) = \frac{1/3 * 3/8}{3/8} = \frac{0.125}{3/8}$$

$$p(\text{female} \mid \text{drew}) = \frac{2/5 * 5/8}{3/8} = \frac{0.250}{3/8}$$

So far we have only considered Bayes Classification when we have one attribute (the “*antennae length*”, or the “*name*”). But we may have many features.

How do we use all the features?

$$p(c_j | d) = \frac{p(d | c_j) p(c_j)}{p(d)}$$

Name	Over 170cm	Eye	Hair length	Sex
Drew	No	Blue	Short	Male
Claudia	Yes	Brown	Long	Female
Drew	No	Blue	Long	Female
Drew	No	Blue	Long	Female
Alberto	Yes	Brown	Short	Male
Karin	No	Blue	Long	Female
Nina	Yes	Brown	Short	Female
Sergio	Yes	Blue	Long	Male

- To simplify the task, **naïve Bayesian classifiers** assume attributes have independent distributions, and thereby estimate

$$p(d|c_j) = p(d_1|c_j) * p(d_2|c_j) * \dots * p(d_n|c_j)$$

↑
The probability of
class c_j generating
instance d , equals....

↑
The probability of class c_j
generating the observed
value for feature 1,
multiplied by..

↑
The probability of class c_j
generating the observed
value for feature 2,
multiplied by..

↑
The probability of class c_j
generating the observed
value for feature n ,
multiplied by..

- To simplify the task, **naïve Bayesian classifiers** assume attributes have independent distributions, and thereby estimate

$$p(d|c_j) = p(d_1|c_j) * p(d_2|c_j) * \dots * p(d_n|c_j)$$

$$p(\text{officer drew}|c_j) = p(\text{over}_{170_{\text{cm}}} = \text{yes}|c_j) * p(\text{eye} = \text{blue}|c_j) * \dots$$



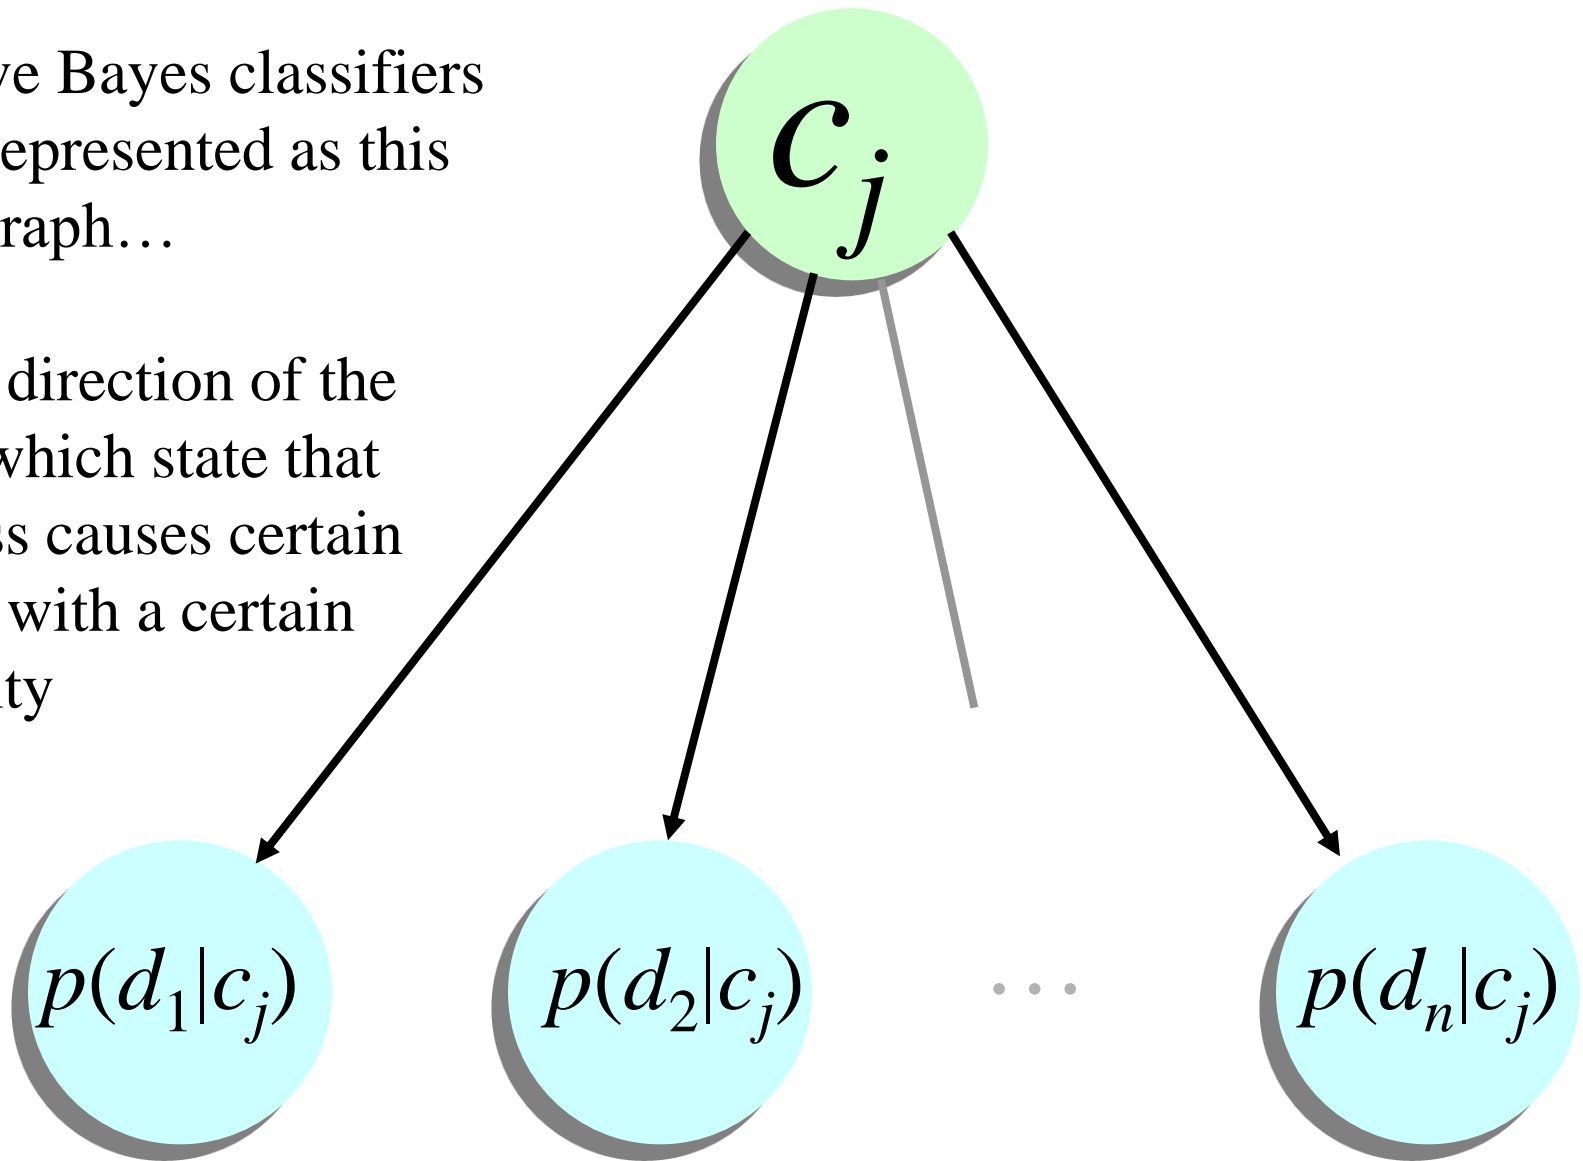
Officer Drew
is blue-eyed,
over 170_{cm}
tall, and has
long hair

$$p(\text{officer drew} | \text{Female}) = 2/5 * 3/5 * \dots$$

$$p(\text{officer drew} | \text{Male}) = 2/3 * 2/3 * \dots$$

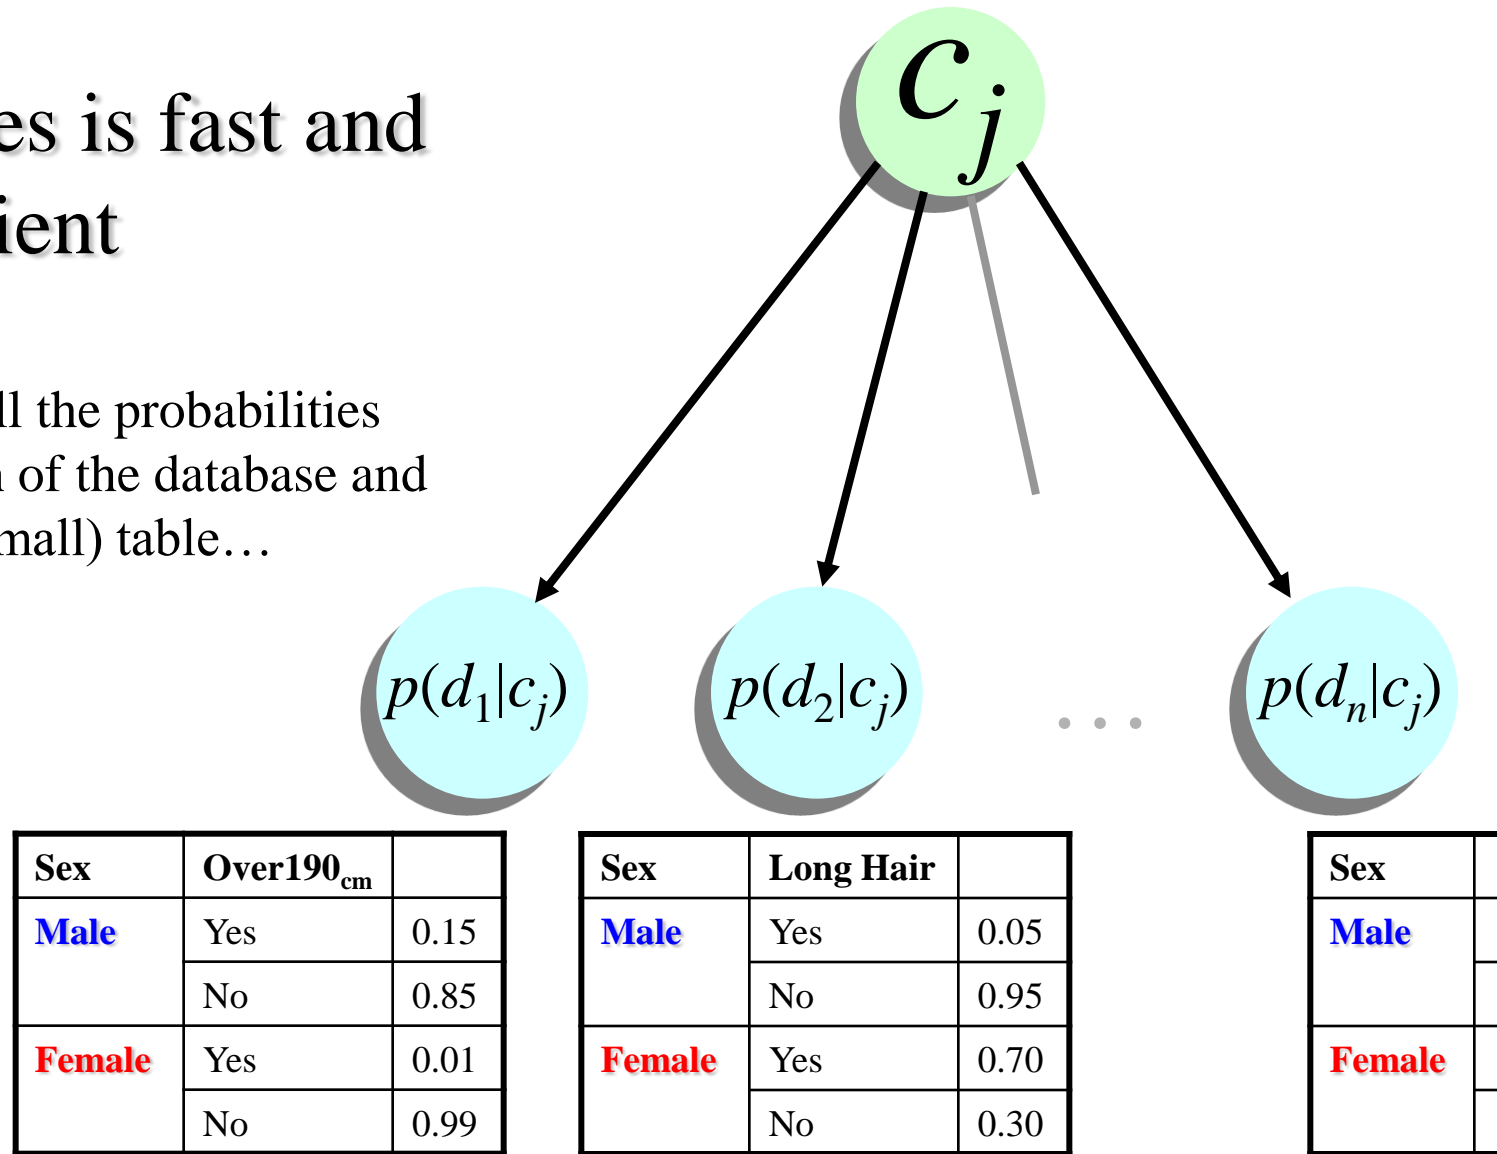
The Naive Bayes classifiers is often represented as this type of graph...

Note the direction of the arrows, which state that each class causes certain features, with a certain probability



Naïve Bayes is fast and space efficient

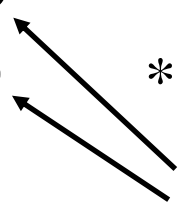
We can look up all the probabilities with a single scan of the database and store them in a (small) table...



Naïve Bayes is NOT sensitive to irrelevant features...

Suppose we are trying to classify a persons sex based on several features, including eye color. (Of course, eye color is completely irrelevant to a persons gender)


$$p(\text{Jessica} | c_j) = p(\text{eye} = \text{brown} | c_j) * p(\text{wears_dress} = \text{yes} | c_j) * \dots$$

$$\begin{array}{llll} p(\text{Jessica} | \text{Female}) = 9,000/10,000 & * & 9,975/10,000 & * \dots \\ p(\text{Jessica} | \text{Male}) = 9,001/10,000 & * & 2/10,000 & * \dots \end{array}$$


Almost the same!

However, this assumes that we have good enough estimates of the probabilities, so the more data the better.

An obvious point. I have used a simple two class problem, and two possible values for each example, for my previous examples. However we can have an arbitrary number of classes, or feature values



Animal	Mass >10 _{kg}	
Cat	Yes	0.15
	No	0.85
Dog	Yes	0.91
	No	0.09
Pig	Yes	0.99
	No	0.01


Animal	Color	
Cat	Black	0.33
	White	0.23
	Brown	0.44
Dog	Black	0.97
	White	0.03
	Brown	0.90
Pig	Black	0.04
	White	0.01
	Brown	0.95

Animal	
Cat	
Dog	
Pig	

Problem!

Naïve Bayes assumes
independence of
features...

Naïve Bayesian Classifier




Sex	Over 6 foot	
Male	Yes	0.15
	No	0.85
Female	Yes	0.01
	No	0.99

Sex	Over 200 pounds	
Male	Yes	0.11
	No	0.80
Female	Yes	0.05
	No	0.95

Solution

Consider the relationships between attributes...

Naïve Bayesian Classifier




Sex	Over 6 foot	
Male	Yes	0.15
	No	0.85
Female	Yes	0.01
	No	0.99

Sex	Over 200 pounds	
Male	Yes and Over 6 foot	0.11
	No and Over 6 foot	0.59
	Yes and NOT Over 6 foot	0.05
	No and NOT Over 6 foot	0.35
Female	Yes and Over 6 foot	0.01

Solution

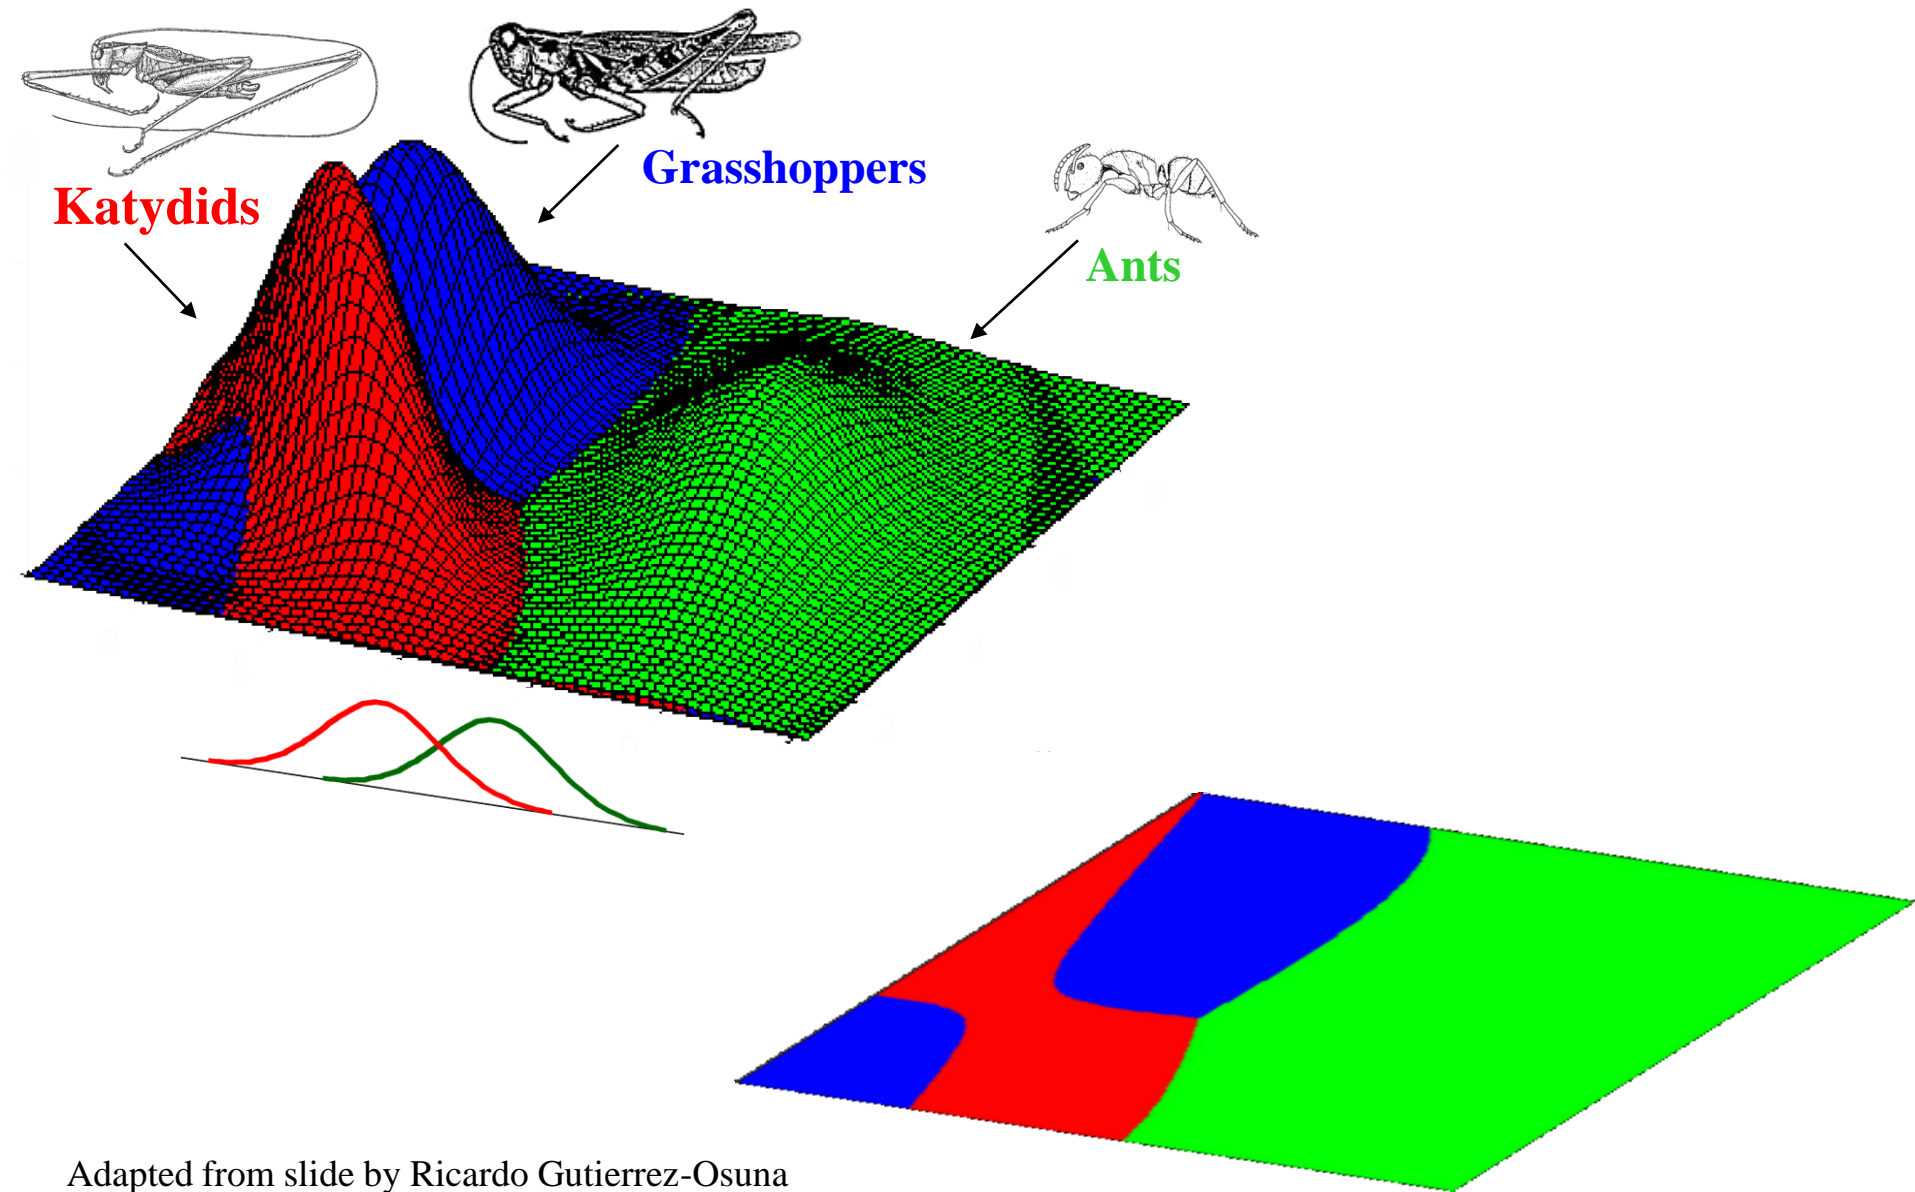
Consider the
relationships between
attributes...



Naïve Bayesian
Classifier






But how do we find the set of connecting arcs??


The Naïve Bayesian Classifier has a piecewise quadratic decision boundary









$$\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$









If I see an insect with a wingbeat frequency of 500, what is it?

$$P(\text{Anopheles} | \text{wingbeat} = 500) = \frac{1}{\sqrt{2\pi} \cdot 30} e^{-\frac{(500-475)^2}{2 \times 30^2}}$$



What is the error rate?

12.2% of the
area under the
pink curve




8.02% of the
area under the
red curve




Can we get more features?

Circadian Features




Suppose I observe an insect with a wingbeat frequency of 420Hz

What is it?




Suppose I observe an insect with a wingbeat frequency of 420Hz at **11:00am**

What is it?



Suppose I observe an insect with a wingbeat frequency of 420 at 11:00am

What is it?






$$(Culex \mid [420Hz, 11:00am]) = \left(\frac{6}{(6 + 6 + 0)} \right) * \left(\frac{2}{(2 + 4 + 3)} \right) = 0.111$$

$$(Anopheles \mid [420Hz, 11:00am]) = \left(\frac{6}{(6 + 6 + 0)} \right) * \left(\frac{4}{(2 + 4 + 3)} \right) = 0.222$$

$$(Aedes \mid [420Hz, 11:00am]) = \left(\frac{0}{(6 + 6 + 0)} \right) * \left(\frac{3}{(2 + 4 + 3)} \right) = 0.000$$

Which of the “Pigeon Problems” can be solved by a decision tree?



Advantages/Disadvantages of Naïve Bayes

- Advantages:
 - Fast to train (single scan). Fast to classify
 - Not sensitive to irrelevant features
 - Handles real and discrete data
 - Handles streaming data well
- Disadvantages:
 - Assumes independence of features